

AI Critique

AI was useful for rapidly translating requirement-based cases into a consistent Playwright structure, but its output was wrong in two different ways, and only the first was easy to see.

The visible errors were mechanical. The browser-matrix runner tried to spawn `npx.cmd` directly on Windows and then reported nine failed runs even though no Playwright process had started; it read the child exit status without checking the spawn error or confirming that report files existed. The OTP assertion used a word-boundary regex, but React rendered adjacent elements with no whitespace between them, so a correct six-digit value failed. I fixed the runner to call the resolved Playwright CLI through Node, switched to non-digit boundaries, and added a validator that opens every embedded `report.json` to check runner identity, timestamp, feature, browser, and totals.

The dangerous error was invisible in exactly the evidence I had trusted. Three fixtures were written so that the case handed the fixture the answer it was about to assert. The clearest was the FR-11 ownership case: it claimed to verify that another user's orders stay hidden, while the mocked endpoint returned only the current user's orders. It passed on all three browsers, and would have passed against a system with no ownership check whatsoever. When I rebuilt it on the mixed data the requirement actually prescribes, it failed immediately and exposed a real defect — the profile table renders another user's order verbatim.

My earlier lesson, that execution evidence must drive acceptance, was incomplete. These tests executed and produced green evidence. The sharper principle is that a passing test is worthless unless I can state what change to the product would make it fail, and the AI never reasons about that unless asked.