

AI Critique – Student Review Required

AI was useful for rapidly translating requirement-based cases into a consistent Playwright structure, but its first output was not execution-ready. The clearest error was the browser-matrix runner: it attempted to spawn `npm.cmd` directly on Windows and then reported nine failed runs even though no Playwright process had started. The AI did not initially distinguish an orchestration failure from product-test failures because it relied on the child exit status without checking the spawn error or confirming report files. I corrected the runner to execute the resolved Playwright CLI with Node and added a validator that opens every embedded `report.json` and checks runner identity, timestamp, feature, browser, and totals.

The AI also used a word-boundary regular expression for the visible OTP. React rendered adjacent elements without whitespace in their combined text, so the correct six-digit value failed the assertion. I replaced word boundaries with non-digit boundaries. I also scoped the back-to-login locator to the main feature area; otherwise the global header login link produced a false pass. Finally, the initial browser plan assumed Firefox and WebKit were available. After installation was unavailable, I verified the installed executables and explicitly used Chromium, Chrome, and Edge while documenting the shared-engine limitation.

The main lesson is that AI can generate a strong automation skeleton, but execution evidence must drive acceptance. Every selector, wait, assertion, subprocess call, and summary needs an independent human oracle. A plausible script is not a valid test until it runs, fails for the intended reason, and produces attributable evidence.