

HW05 - AI-Assisted Performance Testing

1. Student and submission information

Field	Value
Student	NGUYEN THIEN NHA TRAN
Student ID	23127272
Class	23KTPM2
Assignment	HW05-AI
Date	2026-08-17
Tool	Apache JMeter 5.6.3
SUT	EShop REST backend, Node.js + Express + SQLite
Selected workflow	Scenario D - Admin order fulfillment

This report separates measured facts from required student-captured evidence. Raw logs and machine traces in [performance/results/](#) were produced by real local executions. No screenshot, video URL, GitHub issue URL, or endurance value is invented.

2. Scenario selection and source review

Scenario D was selected because it differs from the common customer checkout flow and still covers the three required endpoint groups.

Step	Endpoint	Group	Correlation / assertion
1	<code>POST /api/login</code>	Auth-heavy	Extract JWT; assert token and <code>user.role == admin</code>
2	<code>GET /api/users/me</code>	Auth-heavy	Assert email and admin role
3	<code>GET /api/admin/orders</code>	Read-heavy	Assert CSV order exists with <code>current_status</code>
4	<code>GET /api/products</code>	Read-heavy	Assert non-empty JSON array
5	<code>GET /api/categories</code>	Read-heavy	Assert non-empty JSON array
6	<code>PUT /api/admin/orders/{id}/status</code>	Transactional	Assert business-success message
7	<code>GET /api/admin/orders</code>	Read-heavy verification	Assert the same order now has <code>next_status</code>

Source review found four facts that shaped the plan:

1. The database is dropped and rebuilt whenever `server.js` starts. Restart is therefore the reset procedure.
2. Admin routes authenticate a JWT but do not enforce the admin role. The test adds an explicit role pre-check and records BUG-ADMIN-001.
3. Order transitions are stateful. Every iteration must receive a unique order ID; CSV recycling is disabled.
4. `GET /api/admin/orders` has no pagination and returns every order. With 6,000 seeded rows, the baseline response is about 1.07 MB, making it a meaningful read-heavy operation.

3. Data design and reset contract

`performance/data/admin-orders.csv` has these columns:

```
admin_email,admin_password,order_id,current_status,next_status
```

The seed tool creates 6,000 unique pending orders for Load/Stress/Spike and 7,500 for Endurance. Every measured run follows the same sequence:

1. Verify the process on port 3000 is the EShop `node server.js` process.
2. Restart it, which resets SQLite, lockout counters, and in-memory state.
3. Seed orders through the public checkout API; setup traffic is excluded from JTL.

4. Run preflight: unique IDs, all states pending, valid admin login, explicit role check.

5. Run JMeter headlessly while sampling the backend PID every second.

CSV uses `shareMode.all`, `recycle=false`, and `stopThread=true`. This prevents two virtual users from updating the same mutable order. Valid credentials avoid contaminating the workflow with the account-lockout path.

4. Baseline and workload model

The baseline used one virtual user, five sequential workflows, and no think time. Whole-workflow mean was **193.53 ms**, p95 **246.33 ms**. The two order-list calls dominated: 79.52/77.98 ms mean and approximately 1.07 MB per response. Status update mean was 8.17 ms, p95 11.78 ms.

Plan	Workload	Duration	Intent	Listener/report view
Load	20 VU, ramp 30 s	180 s	Expected admin traffic	Aggregate Report
Stress	25 VU; +25 at 90 s; +50 at 180 s	300 s	Find saturation and degradation	Summary Report
Spike	15 VU baseline; +120 VU at 90 s for 60 s	240 s	Sudden surge plus 90 s recovery	View Results Tree
Endurance	25 VU, ramp 30 s	900 s	Validate sustained threshold and memory trend	Response Time Graph

Think time is non-uniform: 200-500 ms after login, 500-1,200 ms before reading the large order list, 250-600 ms around product/category reads, 600-1,200 ms before the state change, and 300-700 ms before verification. Load ramp-up avoids an accidental spike. Stress adds cohorts. Spike starts its added cohort within one second.

Human-reviewed acceptance criteria:

- Assertion error rate below 1% for a valid Load run.
- Login p95 below the SUT's 700 ms authentication target.
- Product-list p95 below the SUT's 800 ms product-list target.
- Stress is exploratory: saturation is a finding, not an automatic failure.
- Spike recovery passes only if the post-spike baseline returns close to the pre-spike latency/rate with no continuing error burst.
- Endurance threshold requires less than 1% errors, no upward p95 drift, and a bounded memory trend.

5. Test-plan review and fixes

AI draft risk	Why it was wrong/incomplete	Final correction
Reuse one order across VUs	State transitions are irreversible and concurrent updates race	Unique CSV order per iteration; no recycle; preflight
Trust 200 status	Empty/wrong JSON or a stale state could pass	Seven content-aware Groovy assertions and JWT correlation
Treat any JWT as admin	Backend omits role authorization	<code>/api/users/me</code> role assertion; separate security bug
Use one uniform timer	Admin reading/decision steps have different human delays	Six step-specific uniform-random timers
Run GUI under load	GUI listeners consume load-generator resources	CLI run; listeners retained only as the three required report-view types
Accept the first Spike run	The machine slept, so the cohort started after baseline	Preserved under <code>results/discarded</code> ; added process-scoped anti-sleep; reran

The three deliverable plans pass the local structural validator: bounded duration, CSV input, think time, assertions, extractor/correlation, distinct listeners, and exact filename convention.

6. Execution evidence and results

All valid runs used host `TRAN`, JMeter and SUT on the same machine. Resource figures measure the backend PID, not JMeter.

Metric	Load	Stress	Spike
Valid span	179.98 s	299.92 s	239.50 s
Samples	5,533	25,530	10,727

Metric	Load	Stress	Spike
Assertion errors	0	0	0
Error rate	0.000%	0.000%	0.000%
Overall mean	91.79 ms	199.69 ms	495.64 ms
Overall p95	264 ms	622 ms	1,915 ms
Average throughput	30.74 req/s	85.12 req/s	44.79 req/s
Peak 1-second bucket	42 req/s	175 req/s	220 req/s
Peak threads	20	100	135
Backend CPU peak	5.22%	6.20%	5.59%
Backend memory ceiling	167.9 MiB	369.6 MiB	449.1 MiB

6.1 Load

Load passed its correctness criteria. Login p95 was 248 ms and product-list p95 198 ms, both below their SUT targets. The transactional status update was the slowest p95 at 320 ms. No saturation was detected. The 42 req/s peak is not reported as capacity; the best 30-second window under the analysis rule was 33.87 req/s with worst p95 300 ms.

6.2 Stress

The step windows show the cost of concurrency:

Window	Mean rate	Mean bucket p95	Worst bucket p95	Errors
25 VU steady	37.98 req/s	338.95 ms	772 ms	0%
50 VU steady	94.00 req/s	95.33 ms	743 ms	0%
100 VU	126.52 req/s	554.01 ms	1,252 ms	0%

The analyzer detected saturation after second 197: additional threads no longer consistently increased throughput and added latency. The update endpoint had p95 875 ms, higher than every other label. CPU remained low, so the data does not support a CPU-bottleneck claim. Likely causes must be verified with further profiling; SQLite serialization, JSON construction, transfer of the unpaginated order list, and local generator contention are candidates, not proven facts.

6.3 Spike and recovery

Phase	Mean rate	Mean bucket p95	Worst bucket p95	Mean active threads
Pre-spike steady (30-89 s)	30.15 req/s	63.90 ms	121 ms	16.0
Spike (90-149 s)	96.55 req/s	1,512.18 ms	2,923 ms	134.4
Recovery (180-229 s)	28.16 req/s	83.32 ms	194 ms	15.0

The service survived without assertion failures and recovered. Recovery mean rate was 6.6% below pre-spike and mean bucket p95 was 30.4% higher, but the final 20 buckets had p95 values from 57 to 137 ms and all returned to 15 VU. The run demonstrates recovery, with a modest residual latency penalty rather than immediate perfect equivalence.

6.4 Endurance threshold

The 25-VU Endurance run produced **42,094 samples over 899.75 seconds with 0 errors**. After excluding the 30-second ramp-up, it sustained an average **47.55 req/s** with p95 **97 ms**. The best qualifying 30-second window held 50.23 req/s; the 62 req/s one-second peak is not capacity. Therefore, the highest soak-validated stable threshold on this hardware and dataset is **47.55 req/s at 25 VU**, not the higher short Stress burst.

Backend CPU peaked at 5.72%. Memory ceiling was 286.7 MiB. Minute means rose during warm-up (102.4, 152.5, 207.1, then 231.0 MiB) and then stayed roughly 233-240 MiB from minutes 3-14. The first/last tenth means were 119.7/232.9 MiB, but the minute series shows a plateau rather than continuing growth. Response p95 improved from 111.97 ms in the first tenth to 37.91 ms in the last tenth, with no error drift. This run does not show a memory leak.

6.5 Invalid run retained for traceability

The first Spike attempt is stored in [performance/results/discarded/23127272_Spike_20260817_sleep-invalid](#). The machine slept mid-run; its JTL spans 292.61 s, peak threads reached only 32, and the intended 120-thread cohort started after the

baseline ended. It is excluded from every result table above.

7. AI analysis and misinterpretation hunt

The preserved first-pass AI output is [AI docs/AI-initial-analysis.md](#).

AI claim	Verdict	Correct value from raw evidence	Human correction
Stress capacity is 175 req/s	Invalid	175 is one 1-second peak; overall 85.12 req/s; best qualifying 30-second window 136.17 req/s	A momentary peak is not sustained capacity. Use the soak threshold for the final capacity claim.
Spike did not recover because overall p95 is 1,915 ms	Invalid	Recovery window mean bucket p95 83.32 ms; final 20 buckets 57-137 ms	Overall p95 mixes baseline, spike, and recovery and cannot answer recovery.
Load proves a memory leak	Invalid	First/last tenth 66.4/151.1 MiB, but Load includes ramp-up from 1 to 20 VU	A short ramped run cannot prove a leak. Use the 15-minute resource trend.
Every Stress endpoint has p95 622 ms	Invalid	Per-label p95 ranges from 496 ms (users/me) to 875 ms (status update)	Overall percentile is a weighted distribution, not each label's percentile.
CPU is the bottleneck	Unsupported	Backend CPU peak 6.20% on a 20-logical-processor host	Latency degradation with low process CPU needs DB/event-loop/I/O profiling; no bottleneck is proven.

Optimization review


Proposal	Classification	Evidence-based judgment
Paginate admin order list	Feasible	Source uses an unbounded <code>SELECT ... ORDER BY</code> ; baseline responses are ~1.07 MB and the two list calls dominate read cost.
Add index on <code>orders.status</code>	Unfounded for this flow	The measured list does not filter by status; the update locates <code>id</code> , already the primary key.
Increase DB connection pool	Hallucinated	The backend uses one <code>sqlite3.Database</code> ; no configurable pool exists in this architecture.
Enable SQLite WAL	Feasible experiment, not proven fix	WAL may reduce read/write interference, but these logs do not isolate SQLite locking. Benchmark before/after.
Add admin-role middleware	Feasible security fix	Required by FR-12 and source-confirmed, but it does not explain measured latency.

8. AI Critique (200-300 words)

The AI was useful for turning Scenario D into repeatable JMeter plans, but its first analysis was too willing to convert visible summary numbers into conclusions. It called the Stress peak of 175 requests per second “capacity,” although the raw JTL shows that this was only one one-second bucket; average throughput was 85.12 requests per second, and sustained behavior must be checked over a longer window. It also used the Spike run’s overall p95 of 1,915 ms to claim that recovery failed. That statistic mixes three different phases. The recovery-only buckets instead averaged 83.32 ms at p95 and returned to 15 active threads. A third mistake was treating Load memory growth as a leak. The first tenth occurred during ramp-up, so comparing it directly with the fully loaded final tenth confounds concurrency with leakage. The AI also proposed a connection-pool change that does not exist in this SQLite architecture and an index on `status` that the measured query would not use.

These errors happened because a generic summary hides time windows, endpoint labels, source architecture, and test-side events. The machine-sleep incident made this especially clear: JMeter exited successfully and generated a dashboard, yet the raw timestamps and `allThreads` proved that the first Spike attempt was invalid. My main lesson is that AI should generate falsifiable claims, not final verdicts. I must recompute percentiles from the raw JTL, align them with resource timestamps, inspect each endpoint separately, verify suggestions against source code, and preserve discarded runs. AI accelerates the work; the human remains responsible for the evidence contract and interpretation.

9. Continuous performance testing proposal



The gate should compare the same endpoint label, data snapshot, worker class, and steady-state window. A 15% p95 regression threshold is large enough to reduce noise on shared CI but still catch meaningful degradation; it must be calibrated from repeated runs. A second clean-worker run lowers false alarms but doubles cost on suspected regressions. PRs run only a short Load test; Stress runs nightly and Endurance weekly to control compute time. Raw JTL, environment metadata, and resource traces are retained so a regression can be reproduced. Risks remain: shared runners create noise, fixed seed data may not represent production growth, and broad path filters may skip an indirect performance change. Conservative triggers cost more but reduce false negatives.

10. Required student-captured evidence

- [] Four same-frame screenshots: JMeter terminal + Task Manager at peak for Load, Stress, Spike, Endurance.
- [] `dxdiag`/hardware screenshot with hostname **TRAN**.
- [] At least six minutes of Vietnamese narration; add unlisted YouTube URL.
- [] Publish BUG-ADMIN-001 and attach its screenshot; replace GitHub URL.
- [] Review all AI artifacts and set final audit verdicts/signature.
- [] Confirm Scenario D is not duplicated by another group member.

11. References

1. HW05 assignment, version 2.0.
2. Course lecture [S11.2_Performance Testing.pdf](#).
3. EShop repository and backend source: <https://github.com/ttbhanh/eshop-sut>.
4. Apache JMeter 5.6.3 documentation: CLI mode, CSV Data Set Config, assertions, listeners, and dashboard generation.